

Technical SEO Audit & Remediation Plan

Fusion Dental Implant Center — Russian-Language Section

Document purpose: client-facing implementation plan and acceptance criteria. This is not authorization to change the English-language site or shared production infrastructure without regression testing.

Executive Summary

The supplied audit contains several valid findings, but its recommendations should not be implemented as one undifferentiated checklist. The highest-value work is the multilingual/indexation layer, deterministic server output, structured-data accuracy, localization quality, internal crawl paths, and measurable mobile performance. Several other items are hygiene tasks or optional improvements, while a few recommendations are technically incorrect or unnecessarily risky.

The Russian and English versions share one domain, codebase, templates, caching/CDN behavior, analytics, and conversion paths. Every template-level change must therefore be treated as a whole-site release even when the SEO scope is limited to /ru/.

Recommended implementation order

- First: freeze a baseline crawl and release test set for both EN and RU.
- Second: repair hreflang/canonical/indexability relationships and stabilize the server-rendered language output.
- Third: correct structured-data identity and rating inconsistencies, then remove English interface leakage from RU templates and add transparent original/translation handling for patient reviews.
- Fourth: fix high-value RU metadata, broken internal links, redirecting internal links, and mobile Core Web Vitals.
- Fifth: complete medium-priority template hygiene and optional enhancements only after P0/P1 verification.

Key reviewer conclusions

- Accept: bidirectional hreflang remediation, sitemap completeness review, deterministic lang output, broken-link cleanup, direct internal links, interface localization, transparent review-translation UX, structured-data corrections, and targeted performance work.
- Accept with revision: HSTS, headings, image optimization, alt text, titles/H1/meta descriptions, cookie/privacy work, and HTML sitemap recommendations.
- Do not implement as written: nofollow on every external link, rewriting Next.js image optimization URLs solely to make them “SEO-friendly,” forcing every H1 to differ from the title, enforcing a universal 160-character description limit, converting all flagged headings to divs, or using a green PageSpeed score as the acceptance criterion.

Scope, Evidence, and Guardrails

- Primary SEO scope: Russian-language URLs under /ru/.
- Shared-site safeguard: verify EN behavior after every change to routing, metadata generation, templates, schema, caching, image delivery, headers, analytics, or consent logic.
- Evidence base: the contractor's base technical audit, the structured-data audit, the remaining technical SEO issues report, the semantic-core review, prior project discussions, and a current public-page spot-check.
- Confirmed public observation: the RU homepage still contains English review-module labels, CTA text, service descriptions, and interface fragments. English patient-authored reviews are authentic source content and require transparent translation handling rather than silent replacement.
- Not independently re-crawled in full: every URL list attached to the contractor audit. Counts from those reports should be revalidated immediately before implementation.
- No production edit should be marked complete without a before/after URL list, test evidence, and confirmation that EN templates and conversions were not damaged.

Priority Definitions

- P0 — Blocking / immediate: can impair multilingual interpretation, indexation consistency, or safe deployment across the shared site.
- P1 — High: meaningful SEO, trust, crawl, conversion, or structured-data impact; implement after P0 safeguards.
- P2 — Medium: useful quality, accessibility, social-sharing, security, or maintenance improvement; schedule after P0/P1.
- P3 — Reject or defer: not an SEO requirement, incorrectly generalized, low-value relative to cost, or risky if applied sitewide.

P0 — Blocking / Immediate

P0-1. Rebuild reciprocal hreflang clusters and align canonicals

- Evidence: the supporting audit reports 128 indexable RU URLs, but only 8 complete EN↔RU pairs; 120 RU pages reportedly lack a return link from the EN version.
- Required action: maintain one translation-pair registry and generate the same self-reference, alternate-language reference, and agreed x-default on both versions. Never point an untranslated URL to an unrelated homepage.
- Canonical rule: each indexable language page should normally self-canonicalize. Do not canonicalize RU pages to EN merely because the content topics overlap.
- Validation: every hreflang target returns 200, is indexable, does not redirect, matches its canonical, and returns the same reciprocal cluster.

- Client decision: confirm whether the target codes should be en-US / ru-US or generic en / ru, and confirm the x-default destination.
- Owner: developer, with SEO mapping supplied and approved before implementation.
- EN/shared-site risk: critical. A global generator change can invalidate EN annotations or canonicals across the domain.

P0-2. Make server-rendered language output deterministic

- Evidence: the same /ru/ URL has reportedly returned lang="ru" or lang="en" depending on device, user agent, cache, cookies, JavaScript, or rendering path; Google Search Console captured an English value on a Russian page.
- Required action: /ru/ must consistently return <html lang="ru"> in the initial server response; EN pages must return <html lang="en">. Review locale resolution, hydration, edge/CDN cache keys, cookie behavior, and shared layouts.
- Important nuance: Google does not use the HTML lang attribute by itself to detect page language. The priority comes from deterministic rendering, accessibility, browser behavior, and the possibility that the same inconsistency also affects metadata, schema, hreflang, or content.
- Validation: test representative EN and RU URLs with desktop/mobile user agents, clean sessions, different cache states, JavaScript disabled where possible, and Search Console live inspection. HTML source and hydrated DOM must agree.
- EN/shared-site risk: critical. Fix the shared locale mechanism, not only individual RU pages.

P0-3. Establish a single indexability inventory and repair XML sitemap coverage

- Evidence: the supporting audit reports 30 indexable URLs outside the XML sitemap, including 29 RU location pages and one EN gallery page.
- Required action: create one inventory containing final URL, language, template, HTTP status, indexability, canonical, hreflang cluster, sitemap inclusion, and strategic decision.
- Decision gate: add a URL to the sitemap only if it is canonical, indexable, returns 200, and is intended to rank. Otherwise change its indexability or consolidate it; do not use the sitemap to legitimize thin or unapproved location pages.
- Location-page safeguard: confirm which locations are actual offices versus service-area pages. Avoid large sets of near-duplicate city pages without distinct patient value.
- Validation: XML sitemap contains only approved canonical 200 URLs; lastmod reflects substantive changes; sitemap and HTML hreflang logic agree.
- EN/shared-site risk: high because sitemap generation is shared and the reported gap includes an EN URL.

P0-4. Create release gates before template-level implementation

- Required action: capture a pre-change crawl and a representative QA set covering EN/RU homepages, service pages, articles, locations, contact, insurance, case galleries, and 404 pages.
- Required action: deploy through a reviewed GitHub pull request and staging/preview environment where available; document changed templates and generated metadata.
- Release gate: no unexpected changes to EN titles, canonicals, hreflang, schema, internal links, forms, phone links, GTM/GA4 events, or Salesforce lead delivery.
- Rollback requirement: keep the prior production version and configuration available until post-release checks pass.

P1 — High Priority

P1-1. Localize RU templates without misrepresenting original reviews

- Confirmed observation: the public RU homepage still shows English interface labels, CTAs, service descriptions, gallery text, and review-module labels. English patient-authored review text is a separate case and should not automatically be treated as localization leakage.
- Required action: localize all site-authored template strings under `/ru/`, including CTAs, review-module labels, filters, related-content cards, footer elements, accessibility labels, form states, and validation/error messages.
- Authenticity rule: do not replace an English patient review with Russian copy presented as though the patient wrote it in Russian. Keep the reviewer name, rating, review date, location, source, and original English text unchanged.
- Recommended review widget: on RU pages, show a Russian translation with a visible label and provide an accessible “Show original” / “Show translation” toggle. The original English review remains the authoritative source.
- Translation labeling: use “Translated from English” only for a human-reviewed translation; use “Automatically translated from English” when machine translation is shown. Patient names and other identity fields must never be translated.
- Technical implementation: render the review and translation in indexable, server-delivered HTML where feasible; mark the original text with `lang="en"` and the translation with `lang="ru"`; make the toggle keyboard- and screen-reader-accessible and ensure both states work without changing the EN template.
- Data integrity: retain the source review ID and original text, store the translation as a separate field, and refresh or invalidate it if the source review changes. Do not overwrite the source review body.
- Platform option: Places API (New) exposes both localized review text and `originalText`, but returns a limited review set. Before implementation, confirm the current review source/widget, Google Maps and author attribution, API limits, caching rules, and

product terms. If the current widget cannot expose both versions, use a custom accessible component or an approved widget that can.

- Structured-data rule: any Review markup must match the visible review, rating, author, source, and language presentation. A machine translation must not be represented as the unqualified original review text.
- Validation: compare every displayed review against its source; test labels and toggle behavior on mobile and desktop; inspect initial HTML and the hydrated DOM; verify that Google Maps/source attribution remains visible and accurate.
- EN/shared-site risk: high if locale strings or review data are stored in shared components. EN review rendering and copy must remain unchanged.

P1-2. Correct page identity in structured data

- Evidence: the structured-data report states that 297/297 tested URLs repeat a homepage MedicalWebPage identity instead of describing the current URL.
- Required action: generate WebPage or MedicalWebPage per canonical URL with page-specific @id, url, name, inLanguage, isPartOf, and relevant mainEntity/about relationships.
- Required action: use stable @id references for shared Organization, LocalBusiness, WebSite, and Person entities instead of reproducing a large commercial graph on every internal page.
- Validation: service, article, location, contact, legal, collection, and 404 templates each expose only entities supported by visible page content and pass syntax plus semantic review.
- EN/shared-site risk: high. This is a shared graph/template change.

P1-3. Resolve rating and review-count inconsistencies

- Evidence: the schema audit reported 4.9/312 sitewide, while visible public content currently presents different combinations, including 4.7 with 740+ patients and 4.9 with 300+ reviews on the same RU homepage.
- Required action: choose one verifiable source of truth and distinguish reviews, patients, cases, and ratings. Do not use these counts interchangeably.
- Required action: remove sitewide AggregateRating/Review markup until the visible value, count, entity, and source are consistent. Do not expect self-serving LocalBusiness review stars in Google.
- Validation: every displayed and marked-up rating/count is identical, current, attributable, and visible on the page where it is marked up.
- Client decision: identify the approved review platform/data source and the precise meaning of each count.
- EN/shared-site risk: high because rating components and schema appear global.

P1-4. Prioritize titles, H1s, and meta descriptions by page value

- Evidence: the contractor supplied separate URL lists for missing titles, missing H1s, missing descriptions, duplicate H1s, and long descriptions.

- Required action: re-crawl and prioritize indexable RU money pages first: full-arch implants, implant rescue, single/multiple implants, zygomatic implants, bone grafting, consultation, and approved locations.
- Title/H1 rule: both should accurately describe the page and may legitimately be similar. Do not rewrite solely because they match.
- Description rule: write unique, useful snippets for important indexable pages. There is no universal 160-character requirement; optimize for clarity and expected display, not a hard cutoff.
- H1 rule: use a clear primary heading and logical hierarchy. Do not treat every repeated or multiple H1 as an error without inspecting the rendered template.
- Validation: page-to-keyword mapping, unique intent, correct language, no template leakage, and no accidental EN changes.

P1-5. Repair broken internal links and links that pass through redirects

- Required action: replace internal links to 4xx/5xx URLs with the correct live destination or remove them when no valid destination exists.
- Required action: update internal links that currently point to redirecting URLs so they link directly to the final canonical URL; retain necessary redirect rules for external traffic and historical URLs.
- Validation: zero unintended broken internal links in the approved crawl; no avoidable internal redirect chains; navigation, footer, breadcrumbs, related content, and CTA links tested.
- EN/shared-site risk: medium to high when links come from shared navigation or templates.

P1-6. Improve mobile performance using Core Web Vitals, not a color score

- Evidence: the contractor supplied mobile PageSpeed reports for several RU and EN templates but did not provide a root-cause breakdown or field-data baseline.
- Required action: diagnose representative templates using field data where available and lab traces for root causes. Prioritize LCP, INP, and CLS at the 75th percentile, plus conversion usability.
- Likely workstreams to validate: hero/LCP image delivery, responsive image sizes, font loading, unused JavaScript, hydration cost, third-party scripts, caching, lazy loading below the fold, and layout-reserved dimensions.
- Acceptance criterion: material improvement in field Core Web Vitals and no regression to forms, analytics, accessibility, image quality, or EN pages. A green Lighthouse/PageSpeed score is not the project goal.
- EN/shared-site risk: critical because performance code, image components, and scripts are shared.

P1-7. Normalize article and medical structured-data types

- Evidence: the structured-data audit reports widespread use of MedicalScholarlyArticle for ordinary patient guides and inconsistent author/publisher links.
- Required action: use Article or BlogPosting for patient education unless a page is genuinely a scholarly publication. Connect author and publisher to stable Person/Organization @id values.
- Required action: preserve only properties supported by visible content, including headline, description, dates, image, and mainEntityOfPage.
- Validation: each article template uses the correct type; author/publisher entities are not duplicated; medical topics may be represented through about/mentions where appropriate.
- EN/shared-site risk: high because article templates span both languages.

P1-8. Protect analytics and conversions during SEO releases

- Required action: test RU form submissions, click-to-call, consultation CTAs, thank-you behavior, GTM/GA4 events, consent behavior, and Salesforce lead receipt before and after release.
- Required action: preserve tracking identifiers and event names when headings, buttons, templates, cookie logic, or client-side rendering are changed.
- Validation: one documented end-to-end test per primary RU conversion path plus an EN regression test.
- EN/shared-site risk: critical; a technically correct SEO release is unacceptable if lead tracking or conversion elements fail.

P2 — Medium Priority

P2-1. Repair heading semantics template by template

- Required action: inspect page outlines by template and use headings for meaningful sections in a logical hierarchy. Links, logos, form labels, decorative labels, and layout-only text should not become headings.
- Do not bulk-convert every reported H3/H4 to div. Many listed items are legitimate service, FAQ, process, location, or physician section headings.
- Validation: each template has a coherent accessible outline and no heading is removed solely because it is repeated visually.

P2-2. Improve image delivery and alt text selectively

- Required action: compress oversized assets, use modern formats where supported, serve responsive dimensions, and avoid loading unnecessarily large images.
- Do not impose a universal 100 KB limit; acceptable size depends on dimensions, format, quality, placement, and performance impact.

- Alt-text rule: write concise contextual alt text for informative images; use empty alt for decorative images; avoid keyword stuffing and duplicating adjacent captions.
- Validation: reduced transferred bytes and improved LCP/CLS where relevant, with image quality and accessibility preserved.

P2-3. Standardize Open Graph and social metadata

- Evidence: the supporting audit reports missing og:image, og:url, and og:locale across multiple templates.
- Required action: generate page-specific og:title, og:description, og:image, og:url, og:type, and consistent locale values. Use absolute, accessible image URLs.
- Validation: social preview test on representative EN/RU templates; og:url equals the canonical URL; shared templates do not cross languages.

P2-4. Complete supporting schema only where visible and relevant

- Add BreadcrumbList only when visible breadcrumbs exist and match the marked-up hierarchy.
- Use ContactPage, CollectionPage/ItemList, MedicalProcedure, and location business entities only when the visible page content supports them.
- FAQPage must be generated from the same visible Q&A source; markup should not contain hidden or unrelated questions.
- 404/410 pages should retain noindex and avoid the full commercial entity graph.
- Validation: schema tests plus manual semantic review; a passing parser alone is not sufficient.

P2-5. Implement HSTS as a security project, not an SEO shortcut

- Required action: first verify HTTPS availability, redirects, certificates, third-party dependencies, and every required subdomain.
- Enable a safe Strict-Transport-Security policy in stages. Consider preload only after the domain and all required subdomains satisfy preload requirements and the client accepts the operational commitment.
- Do not describe HSTS preload as a direct ranking fix.
- EN/shared-site risk: critical at the infrastructure level; a mistake can make the entire domain or subdomains inaccessible.

P2-6. Review cookie consent and privacy implementation

- Required action: have the client's legal/privacy owner confirm required notices, policy text, consent categories, and regional behavior. SEO should not draft legal compliance rules.
- Technical requirement: consent logic must not block essential page content or break permitted analytics/conversion tracking; tags must respect the selected consent state.

- Validation: banner links, policy pages, consent persistence, tag behavior, and EN/RU language variants tested.

P2-7. Remove obsolete meta keywords as low-risk hygiene

- Removing meta keywords is reasonable cleanup but is not expected to improve Google rankings.
- Implement only when convenient within shared metadata work; do not prioritize above indexation, multilingual, content, or performance issues.

P2-8. Qualify outbound links individually

- Use `rel="sponsored"` for paid links and `rel="ugc"` for user-generated links. Use `nofollow` when the site does not want to endorse or associate with a destination.
- Normal editorial links to trusted professional organizations, licensing bodies, sources, or partners do not need blanket `nofollow`.
- Validation: review by link purpose, not with a global regex or sitewide attribute.

P2-9. Treat an HTML sitemap as optional UX support

- An XML sitemap and crawlable internal architecture are the SEO requirements. A human-readable HTML sitemap is optional.
- Create one only if it materially helps users or exposes otherwise hard-to-reach approved pages; do not use it to compensate for weak navigation or thin location-page architecture.

P3 — Reject or Defer

P3-1. Do not rewrite Next.js image optimization URLs for “SEO-friendly” appearance

- URLs under `/_next/image` are normal delivery endpoints for optimized assets and are not inherently an SEO error.
- Only change the image pipeline if original assets are inaccessible, indexing is required and blocked, caching is poor, or performance tests show a real problem.

P3-2. Do not add `nofollow` to every external link

- Blanket `nofollow` removes useful editorial signals and is not required by Google.
- Classify paid, user-generated, untrusted, and normal editorial links separately.

P3-3. Do not force H1 and title to be different

- Similarity is often natural because both identify the page topic.

- Rewrite only when one element is inaccurate, duplicated across multiple URLs, overly generic, or misaligned with search intent.

P3-4. Do not enforce a universal 160-character description limit

- Google may truncate or rewrite snippets and does not define a fixed ranking threshold at 160 characters.
- Prioritize usefulness, uniqueness, intent alignment, and high-value pages.

P3-5. Do not use a green PageSpeed score as the acceptance criterion

- Use PageSpeed/Lighthouse as diagnostics. Accept work based on field Core Web Vitals, root-cause resolution, conversion usability, and absence of regressions.

P3-6. Do not globally replace reported headings with div elements

- Several reported headings describe real sections and FAQs. A blind replacement can damage accessibility, document structure, and search understanding.
- Correct the component semantics after reviewing each template role.

Implementation Phases and Control Points

Phase 0 — Baseline and decisions

- Export current GSC coverage, performance, Core Web Vitals, and sitemap status for RU and EN.
- Run one full crawl with rendered and source-HTML checks.
- Create the EN↔RU URL-pair registry and decide indexability for every RU location page.
- Confirm approved rating/review sources, language-region codes, x-default destination, and implementation owner.
- Define the staging and rollback process.

Phase 1 — Multilingual and indexation foundation

- Implement reciprocal hreflang/canonical logic.
- Stabilize server-rendered lang and locale-dependent metadata/content.
- Rebuild XML sitemap from the approved indexability inventory.
- Run EN/RU regression tests and Search Console live inspections.

Phase 2 — Shared template and trust fixes

- Remove RU interface localization leakage and implement transparent original/translation handling for patient reviews.

- Correct WebPage identity, global graph scope, ratings, article types, and entity references.
- Repair high-value metadata, broken links, and redirecting internal links.
- Confirm conversions and analytics.

Phase 3 — Performance and supporting quality

- Implement measured Core Web Vitals improvements.
- Optimize responsive images and alt text.
- Repair headings, Open Graph, breadcrumb/template schema, privacy/consent behavior, and security items as approved.

Release Acceptance Checklist

- Representative RU and EN pages return the expected 200/3xx/4xx status.
- Canonicals are correct and language-specific.
- Hreflang clusters are reciprocal, self-referencing, canonical-aligned, and free of redirect/noindex targets.
- Server HTML and hydrated DOM return the correct lang, locale, metadata, and content.
- XML sitemap contains only approved canonical indexable 200 URLs.
- Structured data matches the current page and visible content; ratings and counts use an approved source.
- RU templates contain no unintended English interface or explanatory copy; authentic English reviews retain the original text, and any RU translation is clearly labeled and paired with an accessible original/translation toggle.
- Internal broken links and avoidable redirect hops are removed.
- Forms, phone links, GTM/GA4 events, consent behavior, and Salesforce lead delivery pass end-to-end tests.
- Mobile performance is compared against the baseline using field data where available.
- No unintended EN template, ranking-signal, analytics, or conversion regression is observed.
- GSC live tests and post-deployment crawl results are attached to the completion report.

Client Decisions Required Before Development

- Which RU location pages should be indexable, and which locations are actual offices versus broader service areas?
- Should hreflang target en-US / ru-US or generic en / ru, and where should x-default point?
- What is the authoritative source for ratingValue, reviewCount, patient counts, and other trust metrics? For review copy, approve human-reviewed translation, clearly labeled machine translation, or original-only display.
- Which privacy/cookie requirements and consent platform have been approved by legal/privacy stakeholders?

- Who owns implementation, staging QA, production deployment, rollback, and final SEO acceptance?
- Which changes may touch shared EN templates, and who provides EN-side approval?

Required Contractor Reporting

- For every issue: baseline count, affected URL list, implementation specification, owner, status, deployment reference, and verification result.
- Separate fixed, partially fixed, not fixed, out of scope, rejected, and awaiting client decision.
- Report RU and EN regression results separately.
- Do not mark a broad template task complete based on one manually checked URL.
- Use measurable acceptance criteria from this document rather than statements such as “improved,” “optimized,” or “green zone reached.”

Limitations

This document consolidates and reviews the supplied audits and a public-page spot-check. It does not replace a fresh authenticated crawl, Search Console data, repository review, CDN/cache configuration review, analytics debug session, or legal review. URL counts and issue statuses should be revalidated immediately before implementation and again after deployment.